Master Prompt — The DraftKings Reallocation Call

Cutler Layer Answer Key | Not for distribution to exercise participants

---

The Prompt

You are a senior analytics director at Havas Edge, accountable to the AI Council for any AI-influenced number that moves from an internal model signal to a client-facing dashboard.

Situation: RAMP's incrementality test and the AI-assisted media-mix model disagree on a proposed 14-point budget shift for a DraftKings campaign (pulling spend from linear TV into CTV and paid social). The AI model reads the shift as incremental; RAMP's holdout-based incrementality test does not support the same magnitude of lift. Confidence on the AI recommendation is 71%. The campaign runs in six states, two of which carry active responsible-gambling ad frequency caps and self-exclusion suppression requirements. Two of the paid social data feeds in VantEdge Point are 36 hours stale; the rest are current. The client has requested a reallocation answer within 48 hours.

Do the following, in order:

1. **Data inventory and lineage.** List every VantEdge Point and RAMP field this recommendation depends on. Flag the two stale paid social feeds by name and state, explicitly, whether 36-hour staleness on its own is sufficient grounds to hold the recommendation, independent of the model disagreement.

2. **Attribution reasoning.** Separate correlation from incrementality in the AI model's lift read. State, in one paragraph, exactly where and why the AI's incremental lift estimate diverges from RAMP's holdout test, and quantify the gap.

3. **Compliance routing.** Route this specifically through gaming-vertical rules: identify which of the six states' responsible-gambling ad frequency caps and self-exclusion suppression requirements are implicated by a CTV/paid social increase, and state whether the reallocation as proposed could put any state's cap at risk before a human reviews it.

4. **Failure-mode resolution.** Because the AI recommendation and RAMP disagree at 71% confidence, state the resolution process: what confidence threshold would have to be cleared for this to ship without controls, who owns the decision when confidence falls below that threshold, and what the client-facing message says to DraftKings while that decision is pending.

5. **Human-in-the-loop governance.** Name the specific review owner and escalation path for this exception state. Do not default to “leadership reviews it” — name the role responsible for the sign-off before this reaches a dashboard.

6. **Executive artifact.** Return your answer as a verdict block in this exact format:

```
Attribution Integrity Gate Verdict: [SHIP / SHIP WITH CONTROLS / HOLD /
NO-SHIP]
Why: [one executive sentence]
Evidence: [3-5 bullets covering client trust, compliance, cycle time, failure
mode, release readiness]
Open Risks: [bullets, or “None material”]
Next Gate: [what must happen before this reaches a client dashboard]
```

Format completed by Erwin Maurice McDonald to match the Attribution Integrity Gate's mandatory output block (Havas Edge Engagement Intelligence v1.3). The source document referenced this format but the template did not carry through to the final export.

7. Validation instruction. Before finalizing, review your own answer for anything you assumed rather than verified from the data given. State those assumptions explicitly at the end under an “Assumptions” heading, and do not silently resolve any gap you cannot support with the data provided.

This prompt is the internal scoring reference for the Cutler Layer stress-test exercise. Do not share with exercise participants before or during the exercise — only the scenario brief and the disclosure that submissions will be scored should reach them.