

Havas Edge Director-Level Stress-Test Exercise

The DraftKings Reallocation Call — Use Case Brief

Setup

Greg Cutler, VP Analytics at Havas Edge, brings the intern class together and frames the exercise as a brainstorming session, not a test. Before anyone opens a laptop, he tells them plainly that whatever they submit will be scored against a structured rubric and that a written verdict will come back to him. He does not disclose the rubric weighting, and he does not disclose that the exercise is actually running against the Cutler Layer, the same evaluator he uses to gate real AI Council output before it reaches a client dashboard.

What the interns don't know is that their prompt, not their final written answer, is the thing being graded. Whoever's submission produces the highest Verdict Prompt Score against the Cutler Layer's eight scoring dimensions wins the challenge.

Each intern is asked for four deliverables: a master prompt, an executive summary, their stated assumptions, and a validation plan with a final recommendation.

The Scenario

RAMP just flagged a disagreement. The AI-assisted media-mix model is recommending a 14-point budget shift for a DraftKings campaign, pulling spend out of linear TV and into CTV and paid social, based on a lift signal the model is reading as incremental. RAMP's own holdout-based incrementality test does not support that lift at the same magnitude. Confidence on the AI recommendation sits at 71 percent. The campaign runs in six states, two of which have active responsible-gambling ad frequency caps and self-exclusion suppression requirements. VantEdge Point's most recent data refresh for two of the paid social feeds is 36 hours old; the rest is current. The client wants a reallocation answer inside 48 hours. Build the master prompt you would give an AI system to help you reason through this, before it becomes a dashboard number DraftKings sees.

That is the entire brief handed to participants. It is deliberately unresolved: a model disagreement, a compliance surface most interns will not think to name specifically, a data-freshness gap buried in the details, and a time-pressure line designed to tempt people into skipping governance to move fast.

Why This Use Case Is a Strong Fit for the Havas Edge Skill

The skill exists to answer one governing question: does an AI-influenced attribution or media-mix number move from an internal model signal to a client-trust deliverable safely. This exercise was built to put a participant directly inside that question rather than adjacent to it, and it exercises every structural piece the skill is designed to score.

It puts a real Gate 4 situation in front of a participant

The skill names the Failure-Mode Gate, the disagreement between an AI recommendation and RAMP's existing model or incrementality test, as the single strongest scenario to raise with Greg because it is concrete and almost certainly not yet formally addressed by the AI Council. This scenario is that gate, built out with real numbers: a 14-point shift, a named confidence level, and an explicit contradiction with RAMP.

It forces vertical-specific compliance routing instead of a generic answer

The rubric's Compliance and Vertical Risk dimension rewards routing by vertical, pharma to FDA OPDP, insurance to state ad regs, gaming to state responsible-gambling rules, rather than a single “make sure it’s compliant” line. DraftKings sits in gaming, a vertical the skill's own public-record reference names by client and compliance surface, so a participant either demonstrates that specific knowledge or exposes the exact gap the calibration anchors describe between an operator-level submission and a merely competent one.

It exercises the full eight-dimension rubric, not a subset

Dimension	How the scenario exercises it
Business framing	A named decision, stakeholder, and 48-hour deadline force a clear framing test
Data lineage	The 36-hour-stale paid social feeds create a real lineage question to catch or miss
Attribution / MMM reasoning	A quantified, named disagreement between the AI model and RAMP's holdout test
Compliance and vertical risk	Gaming-vertical responsible-gambling caps across two of six states
Failure-mode handling	71 percent confidence against an unresolved model disagreement, with a live deadline
Human-in-the-loop governance	No stated owner for the exception state; the prompt must supply one
Executive artifact quality	Tests whether the participant asks for a verdict block or a narrative answer
AI-use maturity	The time pressure line tempts an assumption-filling answer instead of a validated one

It matches the skill's own Exercise Build Pattern step for step

- A one-page brief with ambiguous but realistic data, model, and compliance tension
- A disclosure line given to participants before they begin, with the rubric itself withheld
- A request for a master prompt, executive summary, assumptions, and validation plan
- A scored master prompt, not a scored final answer, which is what the Cutler Layer evaluates

It diversifies away from the skill's existing worked example

The skill's reference materials already carry a pharma scenario (Health Media Hub, FDA OPDP, a 12-point MMM disagreement) as the worked SHIPS anchor example. Building this exercise on gaming instead of pharma tests whether a participant's judgment generalizes across verticals rather than pattern-matching to an example they may have seen, which is a stronger signal of real director-level reasoning than a second pharma scenario would be.

It preserves the skill's disclosure requirement by design

The skill is explicit that the mode is disclosed, never covert: participants must know an evaluation and written verdict are coming, even though the rubric weighting stays hidden. This use case is built around Greg stating that up front, which keeps the exercise inside the skill's own operating rules rather than running a hidden evaluation on named individuals.