

The Reallocation Call

Director-Level Stress-Test for AI-Influenced Media-Mix Decisions

AUTHOR Erwin Maurice McDonald	FRAMEWORK Attribution Integrity Gate v1.3	VERTICAL Gaming / Responsible Advertising	DATA Fully synthetic
-------------------------------------	---	---	----------------------------

14 POINT SHIFT	71% AI CONFIDENCE	6 STATES, 2 CAPPED	36h STALE FEEDS	48h CLIENT CLOCK
-------------------	----------------------	-----------------------	--------------------	---------------------

PURPOSE

An AI Council's governing question is whether an AI-influenced attribution or media-mix number can move from an internal model signal to a client-trust deliverable safely. This exercise puts a participant inside that question rather than adjacent to it. What is graded is not the participant's recommendation. It is the master prompt they write to reach one.

SETUP

The Council chair frames the session as a brainstorm, not a test. Before anyone opens a laptop, participants are told plainly that whatever they submit will be scored against a structured rubric and that a written verdict will return to the chair. The rubric weighting is withheld. The evaluator's identity is withheld. Each participant produces four deliverables: a master prompt, an executive summary, stated assumptions, and a validation plan with a final recommendation. The submission that scores highest across the eight dimensions of the Attribution Integrity Gate wins.

THE SCENARIO

PARTICIPANT BRIEF · VERBATIM · THIS IS ALL THEY RECEIVE

The holdout-based incrementality test just flagged a disagreement. The AI-assisted media-mix model is recommending a 14-point budget shift for a national sportsbook client, pulling spend out of linear TV and into CTV and paid social, based on a lift signal the model is reading as incremental. The holdout test does not support that lift at the same magnitude. Confidence on the AI recommendation sits at 71 percent. The campaign runs in six states, two of which have active responsible-advertising frequency caps and self-exclusion suppression requirements. The most recent refresh for two of the paid social feeds in the direct-response database is 36 hours old; the rest is current. The client wants a reallocation answer inside 48 hours. Build the master prompt you would give an AI system to help you reason through this, before it becomes a dashboard number the client sees.

The brief is deliberately unresolved. Four traps are built in: a model disagreement with no stated resolution process, a compliance surface most participants will not name specifically, a data-freshness gap buried in the details, and a time-pressure line designed to tempt people into skipping governance to move fast.

DIMENSION	WHAT THE SCENARIO TESTS	PTS
Business Framing	Names the decision, the stakeholder, and the 48-hour deadline	15
Data Inventory & Lineage	Flags the two stale paid social feeds by name and state; decides whether staleness alone is a hold trigger	15
Attribution / MMM Reasoning	Separates correlation from incrementality; quantifies the AI-versus-holdout gap instead of asserting it	15
Compliance & Vertical Risk	Routes to named state responsible-advertising caps and self-exclusion rules, not "make sure it's compliant"	15
Failure Modes & Disagreement	States the confidence threshold, the decision owner below it, and the client message while pending	15
Human-in-the-Loop Governance	Names the review owner and escalation path; never "leadership reviews it"	10
Executive Artifact Quality	Returns the mandatory verdict block, not a narrative answer	10
AI-Use Maturity	States assumptions explicitly; leaves unsupported gaps open rather than silently filling them	5

Thresholds. 85 and above: Release-Ready (SHIP). 70 to 84: Conditional-Ship, named controls required. Below 70: HOLD or NO-SHIP.

MANDATORY VERDICT BLOCK

Attribution Integrity Gate Verdict: [SHIP / SHIP WITH CONTROLS / HOLD / NO-SHIP]
Why: [one executive sentence]
Evidence: [3-5 bullets: trust, compliance, cycle time, failure mode, readiness]
Open Risks: [bullets, or "None material"]
Next Gate: [what must happen before this reaches a client dashboard]

WHY THIS SCENARIO FITS THE GATE

It is a live failure-mode case. The disagreement between an AI recommendation and the existing incrementality test is the strongest scenario to raise with any AI Council, because it is concrete and rarely addressed formally. This brief is that case with real numbers.

It forces vertical routing. Health to fair-balance and substantiation, insurance to state advertising regulations, gaming to responsible-advertising rules. A participant either shows the specific knowledge or exposes the gap between operator-level and merely competent.

It exercises all eight dimensions. Every rubric line has a concrete trigger in the scenario. Nothing is scored on effort.

It follows the build pattern. One-page brief, realistic tension, disclosure before start, rubric withheld, four deliverables, a scored prompt rather than a scored answer.

It diversifies from the reference case. The gate's worked SHIPS anchor is a health-vertical scenario. Running gaming tests whether judgment generalizes rather than pattern-matches.

It stays disclosed. Participants know an evaluation and written verdict are coming. The rubric is hidden; the fact of scoring is not. Covert evaluation of named people is outside the license.

WHAT PARTICIPANTS RECEIVE

- The scenario paragraph, verbatim
- One disclosure line: your prompt will be scored and a written verdict goes to the chair
- Nothing else: no rubric weighting, no evaluator identity, no answer key

WHAT STAYS PRIVATE

- The reference master prompt (answer key)
- The evaluator package and full scoring logic
- Scored submissions and scorecards
- Any engagement-specific reference material

WHERE IT GOES NEXT

The exercise gates one decision. The AI Council Decision Layer sits above it: every backlog item scored on five signals, hard caps on what reaches the executive, a cut list with scores, evidence tags on every line, and any AI-touched client output routed through this gate before it moves. In the companion synthetic sandbox the scenario above is one backlog row. It scores 5, lands as Decision 1, and the gate returns HOLD.