

Director-Level Stress-Test Exercise

The Reallocation Call: Use Case Brief

Epoch Frameworks LLC | Erwin Maurice McDonald | DACR License v2.6 | All clients and data fictional

Setup

The Council chair brings the participant group together and frames the exercise as a brainstorming session, not a test. Before anyone opens a laptop, they are told plainly that whatever they submit will be scored against a structured rubric and that a written verdict will come back to the chair. The rubric weighting is not disclosed, and the identity of the evaluator is not disclosed.

What participants are not told is that their prompt, not their final written answer, is the thing being graded. The submission that produces the highest score against the eight dimensions of the Attribution Integrity Gate wins the challenge.

Each participant is asked for four deliverables: a master prompt, an executive summary, their stated assumptions, and a validation plan with a final recommendation.

The Scenario

This is the entire brief handed to participants.

The holdout-based incrementality test just flagged a disagreement. The AI-assisted media-mix model is recommending a 14-point budget shift for a national sportsbook client, pulling spend out of linear TV and into CTV and paid social, based on a lift signal the model is reading as incremental. The holdout test does not support that lift at the same magnitude. Confidence on the AI recommendation sits at 71 percent. The campaign runs in six states, two of which have active responsible-advertising frequency caps and self-exclusion suppression requirements. The most recent refresh for two of the paid social feeds in the direct-response database is 36 hours old; the rest is current. The client wants a reallocation answer inside 48 hours. Build the master prompt you would give an AI system to help you reason through this, before it becomes a dashboard number the client sees.

It is deliberately unresolved: a model disagreement, a compliance surface most participants will not think to name specifically, a data-freshness gap buried in the details, and a time-pressure line designed to tempt people into skipping governance to move fast.

Why This Scenario Fits the Gate

The Attribution Integrity Gate exists to answer one governing question: does an AI-influenced attribution or media-mix number move from an internal model signal to a client-trust deliverable safely. This exercise puts a participant directly inside that question rather than adjacent to it, and it exercises every structural piece the gate is designed to score.

It puts a real failure-mode situation in front of a participant. The gate names the disagreement between an AI recommendation and the existing model or incrementality test as the strongest scenario to raise with any AI Council, because it is concrete and almost certainly not yet formally addressed. This scenario is that gate, built out with real numbers: a 14-point shift, a named confidence level, and an explicit contradiction with the holdout test.

It forces vertical-specific compliance routing. The Compliance and Vertical Risk dimension rewards routing by vertical: health to fair-balance and substantiation rules, insurance to state advertising regulations, gaming to state responsible-advertising rules. A participant either demonstrates that specific knowledge or exposes the gap between an operator-level submission and a merely competent one.

It exercises the full eight-dimension rubric, not a subset.

Dimension	How the scenario exercises it
Business framing	A named decision, stakeholder, and 48-hour deadline force a clear framing test
Data lineage	The 36-hour-stale paid social feeds create a real lineage question to catch or miss
Attribution / MMM reasoning	A quantified, named disagreement between the AI model and the holdout test
Compliance and vertical risk	Gaming-vertical responsible-advertising caps across two of six states
Failure-mode handling	71 percent confidence against an unresolved disagreement, with a live deadline
Human-in-the-loop governance	No stated owner for the exception state; the prompt must supply one
Executive artifact quality	Tests whether the participant asks for a verdict block or a narrative answer
AI-use maturity	The time-pressure line tempts an assumption-filling answer instead of a validated one

It matches the exercise build pattern step for step. A one-page brief with ambiguous but realistic data, model, and compliance tension. A disclosure line given before participants begin, with the rubric withheld. A request for a master prompt, executive summary, assumptions, and validation plan. A scored master prompt, not a scored final answer.

It diversifies away from the reference example. The gate's reference materials carry a health-vertical scenario (a 12-point model disagreement under fair-balance rules) as the worked SHIPS anchor. Building this exercise on gaming tests whether a participant's judgment generalizes across verticals rather than pattern-matching to an example they may have seen.

It preserves the disclosure requirement by design. The mode is disclosed, never covert: participants must know an evaluation and written verdict are coming, even though the rubric weighting stays hidden. The exercise is built around the chair stating that up front, which keeps it inside the gate's own operating rules rather than running a hidden evaluation on named individuals.

What Participants Receive

The scenario paragraph above and one disclosure line: your prompt will be scored, and a written verdict goes to the Council chair. No rubric weighting, no evaluator identity, no answer key.

What Stays Private

The reference master prompt (the answer key), the evaluator package, and scored submissions. Publishing the answer key would defeat the exercise.

Epoch Frameworks LLC | Erwin Maurice McDonald | DACR License v2.6. Every client, person, and number in this brief is fictional. Attribution Integrity Gate: McDonald (2026).